

# CBRS Cerebras Systems Inc.

NASDAQ · Technology. Semiconductors (AI accelerators) · Updated 2026-06-09

|                               |                               |                                     |                                                                            |
|-------------------------------|-------------------------------|-------------------------------------|----------------------------------------------------------------------------|
| <b>RECOMMENDATION</b><br>SELL | <b>CURRENT PRICE</b><br>\$238 | <b>12M TARGET</b><br>\$134 (-43.7%) | <b>NEXT EARNINGS</b><br>Est. Q2 FY26 in late July 2026 (UBS models 12/2... |
|-------------------------------|-------------------------------|-------------------------------------|----------------------------------------------------------------------------|

|                               |
|-------------------------------|
| <b>Thesis</b>                 |
| Three bullets — the whole bet |

1. Real franchise, now with a real anchor: the UBS initiation surfaced an OpenAI 750MW prepayment deal (~\$20B+ over its life, expandable to 2GW) plus an AWS disaggregated-inference program. Our prior skills-only model missed both; rebuilt for them, fair value rises from \$51 to about \$134.
2. The gap to spot is execution timing, not technology. Base case risk-adjusts the OpenAI ramp ~12 months behind UBS at ~55-60% utilization (the supply chain must scale roughly 10x) and excludes AWS; that path is worth ~\$119. If OpenAI ramps on schedule and AWS converts, the bull case reaches ~\$255, near UBS's \$300.
3. At \$238 the market already credits the UBS-style on-schedule outcome. Blended fair value \$134 (-43.7%) says pay less or wait for proof of the ramp; the lock-up unlock and OpenAI concentration are the swing risks.

|                                                     |                                  |                    |                                                       |                       |                                   |
|-----------------------------------------------------|----------------------------------|--------------------|-------------------------------------------------------|-----------------------|-----------------------------------|
| <b>Key numbers</b>                                  |                                  |                    |                                                       |                       |                                   |
| Pulled via yfinance-data · validated by Delta audit |                                  |                    |                                                       |                       |                                   |
| LAST CLOSE                                          | \$237.83                         | 52-WEEK RANGE      | \$196.73 to \$386.34                                  | DRAWDOWN FROM HIGH    | -38%                              |
| 1-YEAR RETURN                                       | -23.5%                           | YTD RETURN         | n/a (data gap)                                        | MARKET CAP            | \$52.2B                           |
| ENTERPRISE VALUE                                    | ~\$51.8B                         | NET DEBT           | -\$0.44B (net cash)                                   | FY25 REVENUE          | \$510.0M                          |
| FY25 EBITDA / MARGIN                                | -\$110.8M / -21.7%               | FY25 FCF           | Negative (operating CF near breakeven on prepayments) | FY25 OPERATING INCOME | -\$145.3M (-28.5% margin)         |
| FY25 NET INCOME                                     | +\$237.8M (non -operating items) | FY25 GROSS MARGIN  | 39.0%                                                 | PRICE / SALES (TTM)   | 102x                              |
| FORWARD P/E (FY26E)                                 | ~283x                            | R&D; (FY25)        | \$243.3M (48% of revenue)                             | CASH                  | \$701.7M                          |
| TOTAL DEBT                                          | \$261.8M                         | SHARES OUTSTANDING | ~219.6M                                               | PUBLIC FLOAT          | ~8.67M (~96% locked)              |
| BETA                                                | n/a (data gap)                   | ANALYST CONSENSUS  | UBS: Buy, \$300 PT (8 Jun 2026)                       | MEAN / MEDIAN PT      | \$300 (UBS, sole initiation seen) |

|                        |                |  |  |  |  |
|------------------------|----------------|--|--|--|--|
| BLENDDED DCF<br>TARGET | \$134 (-43.7%) |  |  |  |  |
|------------------------|----------------|--|--|--|--|

Business overview

What the company does, how it makes money, who it sells to

Cerebras Systems designs and sells wafer-scale AI computers built around the Wafer Scale Engine, the largest chip in production. It sells CS-3 systems and turnkey supercomputers, and runs Cerebras Inference, a cloud service that serves large language models at industry-leading speed.

BUSINESS MODEL

Revenue comes from two channels. Hardware (CS-3 systems and full Condor Galaxy supercomputer installations sold or leased to customers like G42 and national labs) and a growing cloud business where customers pay per million tokens or for dedicated capacity on Cerebras-operated infrastructure. The mix is shifting from lumpy system sales toward higher-margin recurring cloud.

SEGMENTS

- Systems & supercomputers (hardware)**  
CS-3 systems and Condor Galaxy clusters. Historically the majority of revenue, anchored by G42. Lower gross margin than cloud.
- Cerebras Inference / cloud**  
Per-token and dedicated-capacity cloud launched August 2024. Faster-growing, higher-margin, the strategic center of the bull case.

TOP CUSTOMERS

|                               |  |                                                                                                                                                               |
|-------------------------------|--|---------------------------------------------------------------------------------------------------------------------------------------------------------------|
| OpenAI                        |  | New anchor per UBS: 750MW prepayment deal (~\$20B+ over its life, expandable to 2GW), structured in three 250MW tranches. The dominant future revenue driver. |
| Amazon / AWS                  |  | Disaggregated-inference collaboration (Trainium prefill + WSE decode); 1,000 CBRS systems in pilot. Upside-only in our base.                                  |
| G42                           |  | Prior anchor (UAE, Condor Galaxy). Majority of 2023 revenue per S-1; now being overtaken by OpenAI.                                                           |
| US national labs & enterprise |  | Argonne, Lawrence Livermore, Sandia; Mistral, AlphaSense, GSK, AstraZeneca, Mayo Clinic, TotalEnergies.                                                       |

GEOGRAPHIC REVENUE

Middle East (UAE, via G42) 0% Americas 0%

Historical financials

P&L; trajectory — revenue, EBITDA, FCF, EPS

| FY   | Revenue (\$B) | Growth  | Gross mgn | EBITDA (\$M) | EBITDA mgn | Net inc (\$M) | FCF (\$M) | EPS     |
|------|---------------|---------|-----------|--------------|------------|---------------|-----------|---------|
| FY22 | 0.02          | —       | 11.7%     | -167         | -678.0%    | -178          | —         | \$-0.81 |
| FY23 | 0.08          | +219.8% | 33.5%     | -123         | -156.7%    | -127          | —         | \$-0.58 |
| FY24 | 0.29          | +268.7% | 42.3%     | -90          | -31.0%     | -482          | —         | \$-2.19 |
| FY25 | 0.51          | +75.7%  | 39.0%     | -111         | -21.7%     | 238           | —         | \$1.08  |

Industry & competitive position

Where the company sits · moat · competitors

AI accelerator spend is the largest capital cycle in tech, and the workload mix is tilting from training toward inference as deployed models proliferate. Inference rewards latency and tokens-per-second, the axis Cerebras optimizes for. A wafer-scale architecture sidesteps the inter-chip networking that caps GPU inference speed, but the dollars attract every well-funded chip vendor.

TAM: \$400B

MOAT

Wafer-scale engineering: single-chip yield, packaging, power and cooling that took nearly a decade to productize, plus software that maps large models onto one chip without manual sharding. The inference speed lead is measurable in benchmarks. The open question is whether speed converts to durable pricing power once rivals and falling token prices compete it away.

COMPETITORS

|                              |                                                                                                                            |
|------------------------------|----------------------------------------------------------------------------------------------------------------------------|
| Nvidia                       | Dominant GPU incumbent; CUDA ecosystem, Blackwell/Rubin roadmap, vastly larger R&D.; The gravity well of the whole market. |
| Groq                         | Inference-speed specialist (LPU); direct competitor on the fast-tokens pitch.                                              |
| SambaNova                    | Reconfigurable dataflow AI systems; competes for enterprise and government deployments.                                    |
| Hyperscaler in-house silicon | AWS Trainium/Inferentia, Google TPU, Meta MTIA, Microsoft Maia; aptive demand designing Cerebras out of their own stacks.  |
| AMD                          | Instinct MI-series accelerators; the credible number-two merchant GPU vendor.                                              |

Bull catalysts

What needs to happen for the long thesis to play out

1 Inference cloud scaling on reasoning-model demand

Cerebras Inference, launched August 2024, runs Llama and other open models at the top of independent speed benchmarks, often several times faster than GPU endpoints. The 2025 wave of reasoning models raised tokens generated per query by an order of magnitude. Cerebras monetizes speed: latency-sensitive buyers pay a premium per million tokens. As inference shifts from cost center to the dominant compute workload, a vendor with a genuine latency edge can capture share GPU clusters struggle to match on the same power budget.

2 OpenAI 750MW prepayment deal as the new anchor

Per UBS, CBRS is the only supplier shipping to OpenAI under a prepayment: a 750MW agreement sized at roughly \$20B+ over its life, structured in three 250MW tranches and expandable to 2GW. There is an OpenAI warrant struck at \$100/share vesting on deployment. This dwarfs the prior G42 base and shifts CBRS toward an inference-cloud model. If it ramps on schedule, UBS models revenue of \$7.1B in 2028 and \$11.2B in 2029; even our risk-adjusted version (a year behind, 55-60% utilization) reaches ~\$8.5B by 2030.

3 AWS disaggregated-inference program (upside)

AWS and CBRS are building a disaggregated solution: Trainium for prefill, Cerebras WSE for decode, AWS's answer to NVDA's GPU-plus-Groq LPX approach. UBS reports 1,000 CBRS systems already in a pilot phase and sizes the opportunity at \$4-10B/yr of hardware before the end of the decade if it scales. It is unsized today, so we keep it out of the base and treat it as bull-case optionality.

4 Wafer-scale moat: WSE-3 and WSE-4 roadmap

The Wafer Scale Engine 3 is a single dinner-plate-sized chip, 4 trillion transistors, 900,000 cores, built to hold an entire large model on one piece of silicon. That removes the networking bottleneck that forces GPU clusters to split models across thousands of devices. Replicating it means solving wafer-scale yield, packaging and power that took Cerebras most of a decade. The WSE-4 roadmap extends the lead.

5 Sovereign AI spend

Gulf governments, and increasingly others, are buying domestic AI compute for strategic reasons rather than pure ROI. Cerebras is already inside that budget via G42 and Condor Galaxy. Sovereign deals are large, multi-year and relatively price-insensitive, and the turnkey supercomputer model fits a buyer that wants capability without building a

GPU networking team.

## 6 US national-lab and government contracts

Argonne, Lawrence Livermore and Sandia run Cerebras systems for scientific computing and large-model training. Government buyers validate the technology, carry long procurement cycles that smooth revenue, and open the door to classified workloads where a US-designed wafer-scale part has a sourcing advantage.

## 7 Gross margin at 39% and climbing on cloud mix

Gross margin reached 39% in 2025, up from near zero in 2022, as the mix moved from one-off system sales toward recurring inference cloud. Cloud carries structurally higher margins. As that mix grows, gross margin has room toward the 50s, the swing factor that turns a loss-making income statement into a cash generator.

## 8 Path to operating profit as revenue outgrows fixed R&D;

R&D; was \$243M in 2025 against \$510M revenue, roughly 48%. That ratio falls fast when revenue doubles and R&D; grows at half the pace. Operating loss was \$145M; at \$1B-plus revenue with a higher gross margin, the model crosses into operating profit. The fixed-cost base is already built.

## 9 Balance sheet de-risked by the IPO

Cash was \$701.7M at year-end 2025 against \$261.8M debt, net cash near \$440M before IPO proceeds. That funds the ramp and WSE-4 without a near-term dilution gun to the head. A clean balance sheet is rare for an AI hardware company still burning cash.

## 10 Scarcity value as the only public wafer-scale pure-play

Investors who want direct exposure to non-GPU AI inference have almost nowhere else to go. Cerebras is the one listed wafer-scale name. That scarcity supports a premium multiple and makes the stock a natural partnership or acquisition target for a hyperscaler or chip incumbent.

### Bear thesis-breakers

Independent bear lane · not visible to the bull author until synthesis

## 1 G42 is the business, in a geopolitical crossfire

The 2024 S-1 showed G42, one UAE customer, at roughly 83-87% of revenue. A company where one Gulf buyer is five-sixths of the top line is a vendor-financed bet on one relationship. US export-control scrutiny of advanced compute sold into the Gulf is rising. If Washington restricts the flow, or G42 builds elsewhere, revenue doesn't dip, it collapses. Watch the 10-K concentration footnote and any 8-K on export licensing.

## 2 102x trailing sales needs impossible numbers

At \$52B market cap on \$510M revenue, the stock trades at 102x sales. To grow into even 15x sales in three years, revenue must reach ~\$3.5B, a 7x off a base that just decelerated from 269% to 76%. Compression alone, from 102x to a still-rich 30x, is a 70% drawdown if revenue merely doubles. The math doesn't agree with the price.

## 3 Lock-up unlock against an 8.67M-share float

Float is ~8.67M shares against ~219.6M outstanding, so ~96% of equity is locked. A thin float manufactures the high price; when the lock-up expires, insiders and pre-IPO funds can sell into a market never tested for real supply. Post-IPO unlocks on hot AI names have repeatedly cut prices 30-50% within weeks. The unlock is the single most predictable catalyst, and it points down.

## 4 The 2025 'profit' is an accounting mirage

Net income was +\$237.8M in 2025 while operating income was -\$145.3M. The gap is non-operating: remeasurement, tax items, one-offs. The real engine still loses money at the operating line on \$510M of revenue. Anyone underwriting 'Cerebras is profitable' off the net income line is reading a number that reverses.

## 5 Growth is already decelerating fast

Revenue grew 220% (2023), 269% (2024), then 76% (2025). The second derivative turned hard. Each step down compresses a 102x multiple violently. The Condor Galaxy buildout is largely a one-time installation cycle; once G42's clusters ship, the comp gets brutal and the hardware line can go flat.

## 6 Single architecture, deep-pocketed competition

Cerebras bets everything on wafer-scale. Nvidia's Blackwell and Rubin, inference specialists Groq and SambaNova, plus hyperscaler in-house silicon (Trainium, TPU, MTIA, Maia) all target the same inference dollars. Cerebras spent

\$243M on R&D; in 2025; Nvidia spends that in days. A benchmark lead is not a moat when rivals own the developer ecosystem (CUDA).

## 7 Inference is becoming a price war

Per-token inference pricing has fallen steeply as Groq, Together, Fireworks and the hyperscalers compete and open models commoditize serving. A 39% gross margin is thin for a 'software-like' story, and a price war pressures the exact line the bull case needs to expand. Speed leadership doesn't guarantee pricing power when buyers route to good-enough endpoints at half the cost.

## 8 Negative book equity, cash burn the IPO only delays

Stockholders equity was -\$579M at year-end 2025 and free cash flow was deeply negative before working-capital swings. The IPO refilled the tank, but a company still losing money operationally while ramping WSE-4 will burn it. If capital markets cool, the next raise is dilutive at a far lower price.

## 9 Customer prepayments flatter the cash flow

G42-style upfront commitments inflate deferred revenue and operating cash flow in ramp years, making the cash statement look better than the unit economics. When the prepayment cycle laps, cash conversion deteriorates even if revenue holds. A classic place where concentrated-customer hardware names look healthier than they are right before a reset.

## 10 No sell-side anchor; a round-trip-prone momentum stock

yfinance returns no consensus target and no analyst count, so price discovery is set by a thin float and momentum, not fundamentals. The stock already ran \$196 to \$386 and back to \$238 inside a year, a 49% peak-to-trough swing. Momentum vehicles unwind when the growth story stalls.

## BEAR CASE IN ONE PARAGRAPH

The bull and bear now hinge on one deal. The OpenAI 750MW agreement is the whole story, and CBRS must scale its supply chain roughly 10x to deliver it. That is the single point of failure UBS itself flags. If deployment slips, or NVDA folds Groq into its roadmap and serves the ultra-fast-inference niche, or the AWS program never sizes, the risk-adjusted ramp we model proves generous and the stock re-rates toward the bear path near \$40. Concentration did not go away; it moved from G42 to OpenAI, just at a larger number. Layer in the IPO lock-up against a ~15% float (UBS) and the post-unlock supply that has cut comparable AI IPOs 30-50%, and the downside is concrete. The 2025 net income headline was an accounting artifact: operating margin was -28.5%, and the company still loses money at the operating line. At \$238 the market prices the OpenAI deal ramping on UBS's schedule. We think a year of slippage and 55-60% utilization is the more likely path, worth about \$134. The bear case does not dispute the wafer-scale technology. It disputes paying today for a 10x execution that has not happened yet.

## Key structural risks

Risks orthogonal to the bear thesis · what breaks the long term, not just the next print

### CUSTOMER OpenAI single-deal dependence

The OpenAI 750MW agreement is now the dominant revenue driver per UBS. A delay, renegotiation, or pivot by OpenAI would gut the model. Concentration moved from G42 to OpenAI at a far larger absolute number, so the single-customer risk is bigger, not smaller.

### EXECUTION 10x supply-chain scale-up

Delivering OpenAI requires scaling CBRS's supply chain roughly 10x from dozens of MW in 2025 to 250MW tranches. UBS calls this the gating factor. Wafer-scale yield, packaging, power and cooling at that volume are unproven, and any stumble slips the whole revenue ramp.

### STRUCTURAL Lock-up overhang on a thin float

Float is ~15% (UBS) with the rest locked. The IPO lock-up expiry is a predictable supply shock that has cut comparable AI IPOs 30-50% within weeks. The quoted price is partly a scarcity artifact until real supply clears.

### EXECUTION NVDA / Groq competitive response

NVDA is folding Groq into its roadmap and has announced a disaggregated LPX solution. If NVDA serves the ultra-fast-inference niche effectively, CBRS's differentiated decode advantage narrows and the premium segment it targets shrinks.

#### REGULATORY

#### Export controls and customer geography

Legacy G42 exposure carries US export-control risk on Gulf compute. Broader AI export policy remains a swing factor for any large compute supplier with international deployments.

#### STRUCTURAL

#### Capital intensity and financing

The OpenAI cloud-heavy early phase is capital-intensive. Customer prepayments fund part of it, but a slip in the ramp pressures cash, and any follow-on raise dilutes. Net income positive in 2025 was driven by non-operating items; the operating line still loses money.

#### MACRO

#### AI capex sentiment reversal

A no-earnings, high-multiple stock priced on out-year revenue is highly sensitive to any rotation out of unprofitable AI names or a broad pullback in AI infrastructure spend.

### Peers

NTM P/E, EV/EBITDA, growth, and YTD / 1Y returns across the peer set

| Ticker | NTM P/E | EV/EBITDA | Rev growth | YTD   | 1Y     |
|--------|---------|-----------|------------|-------|--------|
| CBRS   | 283.0x  | -1000.0x  | +75.7%     | +0.0% | -23.5% |
| NVDA   | 16.5x   | 0.0x      | +85.2%     | +0.0% | +0.0%  |
| AMD    | 37.5x   | 0.0x      | +37.8%     | +0.0% | +0.0%  |
| ARM    | 113.2x  | 0.0x      | +20.1%     | +0.0% | +0.0%  |
| ALAB   | 82.3x   | 0.0x      | +93.4%     | +0.0% | +0.0%  |

**Read:** On price-to-sales, Cerebras at 102x trailing is roughly 5x Nvidia (19.9x) and well above Astera Labs (59x) and Arm (75x), the richest names in the cohort. Nvidia grows faster (85% vs 76%), makes huge margins, and trades at 16.5x forward earnings; Cerebras loses money at the operating line. The peer table doesn't rescue the multiple. It frames how far Cerebras sits above even the most expensive profitable comps, and the closest high-multiple names (ARM, ALAB) are themselves debated as stretched.

### DCF triangulation

Bear / Base / Bull scenarios · WACC × terminal growth × FCF growth path

| Scenario             | WACC   | Terminal g | FCF path   | Implied px   |
|----------------------|--------|------------|------------|--------------|
| Bear                 | 13.0%  | 2.0%       | ×0.45      | \$39.53      |
| Base                 | (base) | (base)     | (base FCF) | \$119        |
| Bull                 | 10.5%  | 4.5%       | ×1.60      | \$255        |
| <b>Current price</b> |        |            |            | <b>\$238</b> |

### Synthesis — three paths

45% × \$119 Base + 25% × \$40 Bear + 20% × \$255 Bull + 10% × \$193 cross-check = \$134. Spot already prices the OpenAI deal ramping on UBS's schedule; we haircut it for the 10x supply-chain scale-up.

**Bull · \$255-300**

**20% probability**

OpenAI ramps on UBS's schedule and AWS scales into a multi-billion hardware stream; revenue approaches UBS's \$11B+ by 2029. Our bull lands near UBS's \$300 PT.

**Base · \$119-150**

**45% probability**



Risk-adjusted OpenAI ramp, a year behind UBS at ~55-60% utilization, AWS optional. Revenue ~\$8.5B by 2030 at hardware-like margins.

Bear · \$40-80

35% probability

OpenAI deployment stalls on the 10x supply-chain scale-up, or NVDA/Groq takes the ultra-fast-inference niche, or the lock-up floods supply. Includes the 10% cross-check anchor.

| Management & capital allocation                                          |                                             |
|--------------------------------------------------------------------------|---------------------------------------------|
| Who runs the company · how they deploy capital · recent insider activity |                                             |
| CEO                                                                      | Andrew Feldman, co-founder & CEO since 2016 |
| CFO                                                                      | Not retrieved via available skills          |

CAPITAL ALLOCATION HISTORY

No dividend and no buyback; this is a cash-burning growth company. Capital has gone into R&D; (\$243M in FY25, 48% of revenue) and manufacturing scale for the WSE-3/CS-3 and the WSE-4 roadmap. The IPO raised primary capital to fund the ramp. With negative operating cash generation, the priority is reaching operating breakeven before the IPO proceeds are consumed, not returning capital.

| Social sentiment                        |
|-----------------------------------------|
| Cross-source signal · Reddit / X / news |

No sentiment skills returned data in this run. Qualitatively, CBRS trades as a retail-and-momentum-driven AI IPO with a thin float; the absence of sell-side coverage means price discovery leans on flow rather than fundamentals. Flagged in data\_gaps.

| Recommendation        |                                                                                                                     |
|-----------------------|---------------------------------------------------------------------------------------------------------------------|
| Action by holder type |                                                                                                                     |
| DIRECTION             | SELL / trim on valuation. The franchise is real; spot already prices the on-schedule OpenAI outcome.                |
| 12M TARGET            | \$134 blended (-43.7%). Base DCF \$119; bull \$255 (near UBS \$300); bear \$40.                                     |
| EXISTING LONGS        | Trim into strength, especially ahead of the IPO lock-up. Let a residual position ride the OpenAI ramp proof points. |
| NEW MONEY             | Wait for evidence the OpenAI 750MW deploys on schedule, or a pullback toward \$120-150, before adding.              |
| HEDGE                 | Long puts dated past the lock-up unlock; the OpenAI-execution binary widens the tails.                              |
| POSITION SIZING       | Small-to-moderate; single-customer (OpenAI) and supply-chain scale-up risk argue against a full-weight position.    |

Catalysts to watch

- OpenAI 750MW deployment cadence: tranche-1 250MW progress and any move toward the 2GW option
- Supply-chain scale-up evidence (the ~10x ramp UBS flags as the gating factor)
- AWS disaggregated-inference program sizing (1,000 CBRS systems in pilot; Trainium-rack attach rate)
- IPO lock-up expiration date and supply into the ~15% float

- NVDA's Groq integration / LPX roadmap as a competitive response in ultra-fast inference
- Quarterly revenue and EBIT-margin trajectory vs UBS's steep 12/27E-12/29E ramp
- OpenAI warrant vesting (\$100 strike, 7-yr) as a deployment milestone signal

| Data gaps                                              |
|--------------------------------------------------------|
| What this cycle could not pull · to be filled next run |

- OpenAI and AWS deal terms are sourced from the UBS initiation (8 Jun 2026), not a primary CBRS filing; our base risk-adjusts UBS's schedule.
- Exact OpenAI tranche timing and pricing not independently verified; modeled at ~55-60% of UBS utilization, ~12 months behind.
- AWS program is unsized (pilot stage); excluded from base, modeled only in the bull scenario.
- FY25 segment split (hardware vs inference cloud) still not in primary data.
- IPO lock-up expiry date not retrievable via available skills; flagged as a key catalyst.
- Beta not available from yfinance; WACC beta 1.30 is an analyst estimate for concentration and post-IPO volatility.